

# LittleBit: Ultra Low-Bit Quantization via Latent Factorization

Methodology Analysis and Detailed Explanation

Kaustubh

November 26, 2025

# Table of Contents

- 1 Introduction
- 2 Methodology Overview
- 3 Step 1: Latent Factorization
- 4 Step 2: Binarization
- 5 Step 3: Multi-Scale Compensation
- 6 Step 4: Dual-SVID Initialization
- 7 Step 5: Residual Compensation
- 8 Training
- 9 Example Walkthrough
- 10 Conclusion

# Introduction to LittleBit

- **Problem:** Deploying LLMs on resource-constrained devices is difficult due to memory and compute costs.
- **Solution:** LittleBit, a novel quantization method targeting sub-1-bit per weight (e.g., 0.1 BPW).
- **Key Idea:** Represents weights using low-rank latent factorization followed by binarization and multi-scale compensation.
- **Performance:** Achieves state-of-the-art results in extreme compression regimes (0.1 - 0.55 BPW).

# Methodology Overview

The LittleBit methodology consists of three main components:

## ① LittleBit Architecture:

- Latent Factorization ( $W \approx UV^\top$ )
- Binarization ( $U_{sign}, V_{sign}$ )
- Multi-Scale Compensation ( $h, g, \ell$ )

## ② Dual-SVID Initialization:

- Smart initialization of parameters using SVD.

## ③ Residual Compensation:

- A secondary path to correct approximation errors.

# Step 1: Latent Factorization

## Concept:

- LLM weight matrices often exhibit low-rank structure.
- Instead of storing the full matrix  $W \in \mathbb{R}^{d_{out} \times d_{in}}$ , we approximate it as the product of two smaller matrices.

## Equation:

$$W \approx UV^\top$$

where:

- $U \in \mathbb{R}^{d_{out} \times r}$
- $V \in \mathbb{R}^{d_{in} \times r}$
- $r \ll \min(d_{out}, d_{in})$  is the latent rank.

## Step 2: Binarization

### Concept:

- To achieve extreme compression, the factors  $U$  and  $V$  are binarized to  $\pm 1$ .
- This reduces storage from 16-bit FP to 1-bit per element.

### Equation:

$$U_{sign} = \text{sign}(U), \quad V_{sign} = \text{sign}(V)$$

Values in  $U_{sign}, V_{sign} \in \{-1, +1\}$ .

## Step 3: Multi-Scale Compensation

### Concept:

- To recover the lost magnitude information, LittleBit introduces learnable FP16 scales.
- Scales are applied across three dimensions: Row, Column, and Latent.

### Scales:

- **Row Scale**  $h \in \mathbb{R}_{d_{out}}$ : Captures output channel magnitude.
- **Column Scale**  $g \in \mathbb{R}_{d_{in}}$ : Captures input channel magnitude.
- **Latent Scale**  $\ell \in \mathbb{R}_r$ : Captures the importance of each latent rank.

### Reconstruction:

$$\hat{W}_{pri} = \text{diag}(h) \cdot U_{sign} \cdot \text{diag}(\ell) \cdot V_{sign}^\top \cdot \text{diag}(g)$$

## Step 4: Dual-SVID Initialization

**Problem:** Random initialization of  $U, V, h, g, \ell$  leads to unstable training.

**Solution:** Dual-SVID.

- ① Perform SVD on  $W$ :  $W \approx U'\Sigma V'^\top$ . Absorb  $\Sigma$  into  $U', V'$ .
- ② Initialize binary factors:  $U_{sign,0} = \text{sign}(U')$ .
- ③ Initialize scales by approximating the *magnitudes*  $|U'|$  and  $|V'|$ .

**Magnitude Decomposition:**

$$|U'| \approx h_0(\ell_{u,0})^\top$$

We use a rank-1 approximation on the magnitude matrix  $|U'|$  to separate the row-wise scale ( $h_0$ ) from the latent-wise scale ( $\ell_{u,0}$ ).

Final latent scale  $\ell_0 = \ell_{u,0} \odot \ell_{v,0}$ .

## Step 5: Residual Compensation

### Concept:

- A single low-rank binary approximation might miss details.
- Add a second "Residual" path to model the error.

### Architecture:

$$\hat{W} = \hat{W}_{pri} + \hat{W}_{res}$$

- $\hat{W}_{pri}$ : Primary approximation (learns main structure).
- $\hat{W}_{res}$ : Residual approximation (learns the error  $W - \hat{W}_{pri}$ ).

Both paths use the same factorized structure but with independent parameters. This effectively doubles the rank but allows for specialized error correction without increasing the storage format complexity (just more parameters).

# Training Strategy

## Quantization-Aware Training (QAT):

- The model is trained end-to-end.
- **Knowledge Distillation:** Uses the original FP16 model as a teacher.
- **Loss Function:**

$$L_{QAT} = L_{out}(KL) + \lambda L_{inter}(MSE)$$

- **Gradient Approximation:** Since  $\text{sign}(x)$  is non-differentiable, **SmoothSign** is used for backpropagation:

$$\frac{\partial \text{sign}(x)}{\partial x} \approx \frac{\partial \tanh(kx)}{\partial x}$$

with  $k = 100$ .

## Example: Input Matrix

Consider a  $4 \times 4$  matrix  $W$ :

$$W = \begin{bmatrix} 1.50 & -0.80 & 0.20 & -1.20 \\ -0.50 & 1.20 & -0.90 & 0.50 \\ 0.80 & -0.20 & 1.50 & -0.50 \\ -1.20 & 0.50 & -0.30 & 1.00 \end{bmatrix}$$

## Example: Step 1 - Latent Factorization (Part 1)

1. Perform SVD on  $W$ :  $W = U\Sigma V^\top$

$$U_{full} = \begin{bmatrix} -0.61 & -0.52 & -0.59 & -0.07 \\ 0.44 & -0.37 & -0.19 & -0.79 \\ -0.45 & 0.70 & -0.54 & -0.13 \\ 0.49 & 0.33 & -0.56 & 0.59 \end{bmatrix}$$

$$\Sigma_{full} = \begin{bmatrix} 3.26 & 0.00 & 0.00 & 0.00 \\ 0.00 & 1.32 & 0.00 & 0.00 \\ 0.00 & 0.00 & 0.83 & 0.00 \\ 0.00 & 0.00 & 0.00 & 0.05 \end{bmatrix}$$

$$V_{full}^\top = \begin{bmatrix} -0.64 & 0.41 & -0.41 & 0.51 \\ -0.32 & -0.00 & 0.89 & 0.32 \\ -0.37 & -0.90 & -0.18 & 0.12 \\ 0.59 & -0.13 & -0.07 & 0.79 \end{bmatrix}$$

## Example: Step 1 - Latent Factorization (Part 2)

2. Truncate to rank  $r = 2$ . Keep top 2 singular values.

$$U_r = \begin{bmatrix} -0.61 & -0.52 \\ 0.44 & -0.37 \\ -0.45 & 0.70 \\ 0.49 & 0.33 \end{bmatrix}, \quad \Sigma_r = \begin{bmatrix} 3.26 & 0.00 \\ 0.00 & 1.32 \end{bmatrix}$$

$$V_r^\top = \begin{bmatrix} -0.64 & 0.41 & -0.41 & 0.51 \\ -0.32 & -0.00 & 0.89 & 0.32 \end{bmatrix}$$

## Example: Step 1 - Latent Factorization (Part 3)

3. Calculate  $\sqrt{\Sigma_r}$  and absorb into factors.

$$\sqrt{\Sigma_r} = \begin{bmatrix} \sqrt{3.26} & 0 \\ 0 & \sqrt{1.32} \end{bmatrix} = \begin{bmatrix} 1.81 & 0.00 \\ 0.00 & 1.15 \end{bmatrix}$$

Calculate  $U' = U_r \sqrt{\Sigma_r}$ :

$$U' = \begin{bmatrix} -0.61 & -0.52 \\ 0.44 & -0.37 \\ -0.45 & 0.70 \\ 0.49 & 0.33 \end{bmatrix} \begin{bmatrix} 1.81 & 0.00 \\ 0.00 & 1.15 \end{bmatrix} = \begin{bmatrix} -1.09 & -0.59 \\ 0.79 & -0.42 \\ -0.81 & 0.80 \\ 0.89 & 0.38 \end{bmatrix}$$

## Example: Step 1 - Latent Factorization (Part 4)

Calculate  $V' = V_r \sqrt{\Sigma_r}$  (shown as  $V'^\top$ ):

$$V'^\top = \sqrt{\Sigma_r} V_r^\top = \begin{bmatrix} 1.81 & 0.00 \\ 0.00 & 1.15 \end{bmatrix} \begin{bmatrix} -0.64 & 0.41 & -0.41 & 0.51 \\ -0.32 & -0.00 & 0.89 & 0.32 \end{bmatrix}$$

$$V'^\top = \begin{bmatrix} -1.15 & 0.75 & -0.74 & 0.92 \\ -0.37 & -0.00 & 1.03 & 0.36 \end{bmatrix}$$

## Example: Step 2 - Binarization

$$U_{sign} = \begin{bmatrix} -1 & -1 \\ 1 & -1 \\ -1 & 1 \\ 1 & 1 \end{bmatrix}, \quad V_{sign}^\top = \begin{bmatrix} -1 & 1 & -1 & 1 \\ -1 & -1 & 1 & 1 \end{bmatrix}$$

## Example: Step 3 - Multi-Scale Compensation (Part 1)

1. Compute Magnitude Matrices  $|U'|$  and  $|V'|$  (element-wise absolute value):

$$|U'| = \begin{bmatrix} 1.09 & 0.59 \\ 0.79 & 0.42 \\ 0.81 & 0.80 \\ 0.89 & 0.38 \end{bmatrix}, \quad |V'|^\top = \begin{bmatrix} 1.15 & 0.75 & 0.74 & 0.92 \\ 0.37 & 0.00 & 1.03 & 0.36 \end{bmatrix}$$

## Example: Step 3 - Multi-Scale Compensation (Part 2)

2. SVD of Magnitude Matrix  $|U'|$ :

$$U_{mag} = \begin{bmatrix} -0.59 & -0.19 & -0.78 & -0.07 \\ -0.42 & -0.07 & 0.27 & 0.86 \\ -0.52 & 0.84 & -0.09 & -0.12 \\ -0.45 & -0.50 & 0.56 & -0.49 \end{bmatrix}$$

$$\Sigma_{mag} = \begin{bmatrix} 2.12 & 0.00 \\ 0.00 & 0.31 \\ 0.00 & 0.00 \\ 0.00 & 0.00 \end{bmatrix}, \quad V_{mag}^\top = \begin{bmatrix} -0.85 & -0.53 \\ -0.53 & 0.85 \end{bmatrix}$$

## Example: Step 3 - Multi-Scale Compensation (Part 3)

3. Truncate to Rank-1 (Keep top singular value  $\sigma_1$  and vectors  $u_1, v_1$ ):

$$u_1 = \begin{bmatrix} -0.59 \\ -0.42 \\ -0.52 \\ -0.45 \end{bmatrix}, \quad \sigma_1 = 2.12, \quad v_1^\top = [-0.85 \quad -0.53]$$

4. Calculate  $\sqrt{\sigma_1} = \sqrt{2.12} = 1.46$ .

## Example: Step 3 - Multi-Scale Compensation (Part 4)

5. Calculate  $h_0$  and  $\ell_{u,0}$  by distributing  $\sqrt{\sigma_1}$ :

$$h_0 = u_1 \sqrt{\sigma_1} = \begin{bmatrix} -0.59 \\ -0.42 \\ -0.52 \\ -0.45 \end{bmatrix} \cdot 1.46 = \begin{bmatrix} -0.85 \\ -0.62 \\ -0.76 \\ -0.65 \end{bmatrix}$$

$$\ell_{u,0}^\top = \sqrt{\sigma_1} v_1^\top = 1.46 \cdot [-0.85 \quad -0.53] = [-1.24 \quad -0.77]$$

## Example: Step 3 - Multi-Scale Compensation (Part 5)

6. Similarly for  $|V'|$ , we obtain  $g_0$  and  $\ell_{v,0}$ :

$$g_0 = \begin{bmatrix} -0.83 \\ -0.46 \\ -0.80 \\ -0.69 \end{bmatrix}, \quad \ell_{v,0}^\top = [-1.25 \quad -0.68]$$

7. Calculate final latent scale  $\ell$  (element-wise product):

$$\ell = \ell_{u,0} \odot \ell_{v,0} = \begin{bmatrix} -1.24 \\ -0.77 \end{bmatrix} \odot \begin{bmatrix} -1.25 \\ -0.68 \end{bmatrix} = \begin{bmatrix} 1.54 \\ 0.53 \end{bmatrix}$$

## Example: Step 4 - Primary Approximation

1. Scale Latent Columns:  $A = U_{sign} \text{diag}(\ell)$

$$A = \begin{bmatrix} -1.54 & -0.53 \\ 1.54 & -0.53 \\ -1.54 & 0.53 \\ 1.54 & 0.53 \end{bmatrix}$$

2. Multiply by Binary  $V$ :  $B = AV_{sign}^\top$

$$B = \begin{bmatrix} 2.07 & -1.01 & 1.01 & -2.07 \\ -1.01 & 2.07 & -2.07 & 1.01 \\ 1.01 & -2.07 & 2.07 & -1.01 \\ -2.07 & 1.01 & -1.01 & 2.07 \end{bmatrix}$$

## Example: Step 4 - Primary Approximation (Cont.)

3. Apply Row Scale  $h$ :  $C = \text{diag}(h)B$

$$C = \begin{bmatrix} -1.76 & 0.86 & -0.86 & 1.76 \\ 0.63 & -1.28 & 1.28 & -0.63 \\ -0.77 & 1.58 & -1.58 & 0.77 \\ 1.35 & -0.66 & 0.66 & -1.35 \end{bmatrix}$$

4. Apply Column Scale  $g$ :  $\hat{W}_{pri} = C\text{diag}(g)$

$$\hat{W}_{pri} = \begin{bmatrix} 1.47 & -0.40 & 0.70 & -1.22 \\ -0.52 & 0.59 & -1.02 & 0.43 \\ 0.65 & -0.73 & 1.26 & -0.53 \\ -1.13 & 0.31 & -0.53 & 0.93 \end{bmatrix}$$

## Example: Step 5 - Residual Approximation

Calculate residual target:  $W_{res\_target} = W - \hat{W}_{pri}$

$$W_{res\_target} = \begin{bmatrix} 0.03 & -0.40 & -0.50 & 0.02 \\ 0.02 & 0.61 & 0.12 & 0.07 \\ 0.15 & 0.53 & 0.24 & 0.03 \\ -0.07 & 0.19 & 0.23 & 0.07 \end{bmatrix}$$

We apply the same LittleBit process (SVD, Binarize, Scale) to  $W_{res\_target}$  to get  $\hat{W}_{res}$ :

$$\hat{W}_{res} = \begin{bmatrix} -0.05 & -0.35 & -0.48 & -0.03 \\ 0.09 & 0.60 & 0.25 & 0.06 \\ 0.07 & 0.53 & 0.22 & 0.04 \\ 0.03 & 0.17 & 0.23 & 0.02 \end{bmatrix}$$

## Example: Step 6 - Final Effective Weight

The final weight is the sum of primary and residual approximations:

$$\hat{W} = \hat{W}_{pri} + \hat{W}_{res}$$

Original  $W$ :

$$\begin{bmatrix} 1.50 & -0.80 & 0.20 & -1.20 \\ -0.50 & 1.20 & -0.90 & 0.50 \\ 0.80 & -0.20 & 1.50 & -0.50 \\ -1.20 & 0.50 & -0.30 & 1.00 \end{bmatrix}$$

Reconstructed  $\hat{W}$ :

$$\begin{bmatrix} 1.42 & -0.75 & 0.22 & -1.25 \\ -0.43 & 1.19 & -0.77 & 0.49 \\ 0.72 & -0.20 & 1.48 & -0.49 \\ -1.10 & 0.48 & -0.30 & 0.95 \end{bmatrix}$$

## Example: Step 7 - Training Loss Calculation

Let random input  $X$  (batch size 2, seq len 1, dim 4):

$$X = \begin{bmatrix} 0.50 & -1.00 & 0.20 & 0.80 \\ -0.20 & 0.50 & 1.00 & -0.50 \end{bmatrix}$$

1. Teacher Output:  $Y_{teacher} = XW^\top$

$$Y_{teacher} = \begin{bmatrix} 0.63 & -1.23 & 0.50 & -0.36 \\ 0.10 & -0.45 & 1.49 & -0.31 \end{bmatrix}$$

2. Student Output:  $Y_{student} = X\hat{W}^\top$

$$Y_{student} = \begin{bmatrix} 0.50 & -1.17 & 0.47 & -0.33 \\ 0.18 & -0.33 & 1.48 & -0.31 \end{bmatrix}$$

## Example: Step 7 - Training Loss Calculation (Cont.)

3. Calculate MSE Loss:  $L = \frac{1}{N} \sum (Y_{teacher} - Y_{student})^2$

$$\text{Diff} = \begin{bmatrix} 0.13 & -0.06 & 0.03 & -0.03 \\ -0.08 & -0.12 & 0.01 & 0.00 \end{bmatrix}$$

$$\text{MSE} = 0.0052$$

# Conclusion

- LittleBit enables extreme LLM compression (down to 0.1 BPW).
- It combines latent factorization, binarization, and multi-scale compensation.
- Dual-SVID initialization and Residual Compensation are critical for performance.
- Results show superior performance over existing methods like STBLLM in sub-1-bit regimes.